

# Convolutional Neural Networks

# Overview

Last time:

- Multilayer perceptron
- Backpropagation

Today:

- Review of MLPs/Backpropagation
- CNNs: neural networks for images

Note

- There are no lecture notes for today's material

# Notation

$$z^{(\ell)} = z^{(\ell)}(x)$$

$$a^{(\ell)} = a^{(\ell)}(x)$$

- Input layer

$$z^{(0)} = x \in \mathbb{R}^d$$

- Hidden layers:  $1 \leq \ell < L$

$$a^{(\ell)} = W^{(\ell)} z^{(\ell-1)}$$

$$\text{where } W^{(\ell)} = [w_{ij}^{(\ell)}]$$

$$z^{(\ell)} = \sigma(a^{(\ell)}) \text{ applied elementwise}$$

where  $d_\ell = \#$  of nodes in layer  $\ell$   
 $d_\ell \times d_{\ell-1}$

- If bias desired, prepend a 1 to any  $z^{(\ell)}$ ,  $0 \leq \ell < L$ , and add column of weights to beginning of  $W^{(\ell+1)}$

- Output layer

$$f(x) = a^{(L)}$$


followed by *identity activation*

- Evaluation of the output from the input is called

*forward propagation*

$$x = \begin{bmatrix} x_1 \\ \vdots \\ x_5 \end{bmatrix} = z^{(0)}$$

# Illustration of Notation



# Backprop

Forward pass:

Using current weights  $\theta$  compute  $f(\mathbf{x}_n)$   
and store intermediate values  $a_{ni}^{(\ell)}, z_{nj}^{(\ell)}$

Initialize backward pass:

For  $i = 1$  to  $d_L$   
  Compute  $\delta_{ni}^{(L)}$

End

Backward pass:

For  $\ell = L - 1$  downto 1

  For  $i = 1$  to  $d_\ell$   
     $\delta_{ni}^{(\ell)} = \sum_k \delta_{nk}^{(\ell+1)} w_{ki}^{(\ell+1)} \sigma'(a_{ni}^{(\ell)})$

  For  $j = 1$  to  $d_{\ell-1}$

$$\frac{\partial R_n(\theta)}{\partial w_{ij}^{(\ell)}} \longleftarrow \delta_{ni}^{(\ell)} z_{nj}^{(\ell-1)}$$

$$w_{ij}^{(\ell)} \longleftarrow w_{ij}^{(\ell)} - \eta \frac{\partial R_n(\theta)}{\partial w_{ij}^{(\ell)}}$$

  End

End

End

# Poll

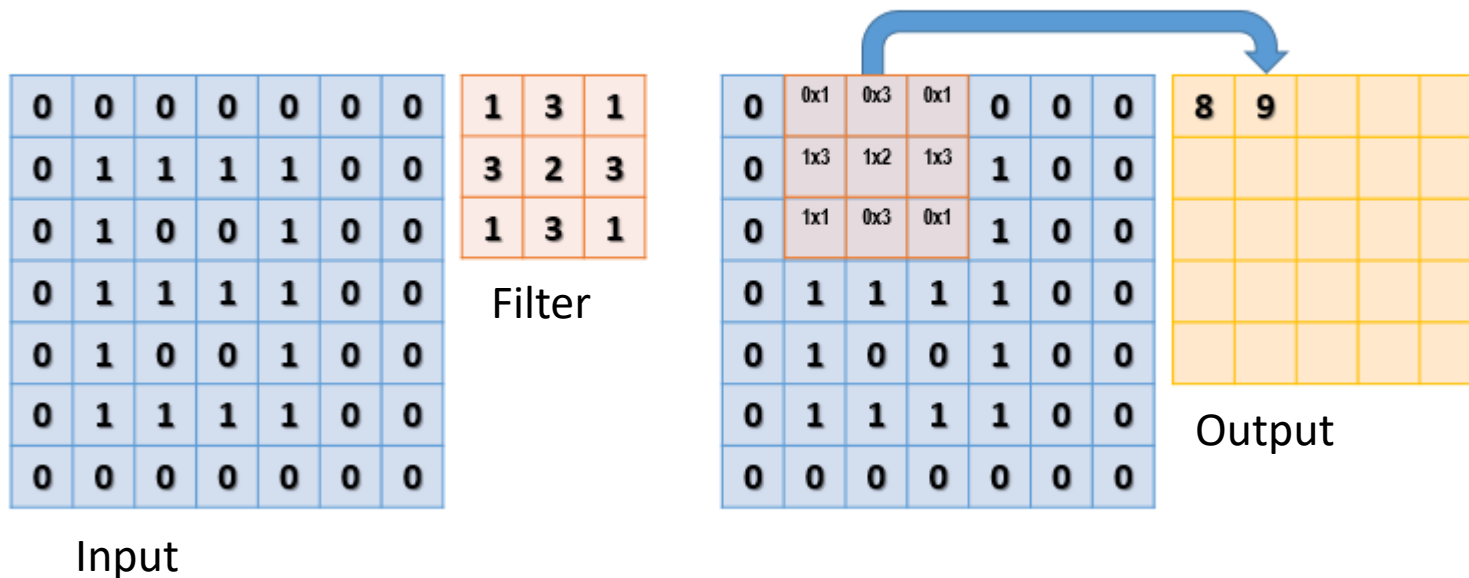
For an MLP with ReLU activation, a subgradient is guaranteed to exist at every iteration of backprop.

(A) True

(B) False

# Convolutions

- Basically a sliding window
- Figure assumes a *stride* (shift increment) of 1, but larger strides are possible
- The filter is also called a



# Poll

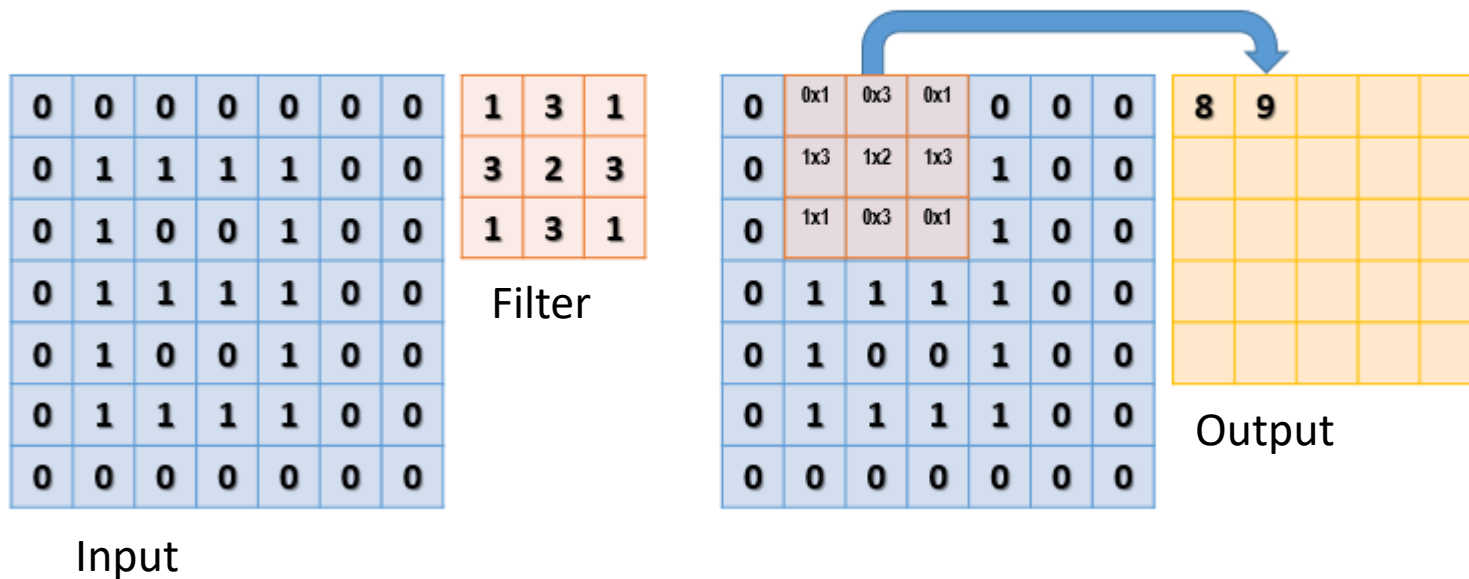
What would be the output size with a stride of 2?

(A)  $2 \times 2$

(B)  $3 \times 3$

(C)  $4 \times 4$

(D)  $5 \times 5$





# Filters are Feature Extractors




Figure: <https://developer.nvidia.com/discover/convolution>

# Padding

0	0	0	0	0	0	0
0	2	4	9	1	4	0
0	2	1	4	4	6	0
0	1	1	2	9	2	0
0	7	3	5	1	3	0
0	2	3	4	8	5	0
0	0	0	0	0	0	0

Image



1	2	3
-4	7	4
2	-5	1

Filter /  
Kernel

=

21	59	37	-19	2
30	51	66	20	43
-14	31	49	101	-19
59	15	53	-2	21
49	57	64	76	10


Feature

# Demo

<https://cs231n.github.io/convolutional-networks/>

# Channels


- Color images have three channels.
- When there are multiple channels, the filter is actually three dimensional, where the depth of the filter is the number of channels.
- The convolution is still *two-dimensional*. Hence the output of the convolution is still two-dimensional.



<https://ai.stackexchange.com/questions/5769/in-a-cnn-does-each-new-filter-have-different-weights-for-each-input-channel-or>

# Convolutional Layers

- The initial layer in a CNN is the input image.
- A convolutional layer is a hidden layer formed by applying several convolutions (each with its own filter) to the previous layer.
- Filter coefficients are the weights to be learned
- # of weights into a layer = (filter  $H \times W \times D$ )  $\times$  (# of filters)



# Conv. Layers are not Fully Connected and Weights are Shared

$v_1$  0

$v_2$  0

$v_3$  0

$v_4$  0

$v_5$  0

$v_6$  0

$v_7$  0

$v_8$  0

$v_9$  0

$v_1$	$v_2$	$v_3$
$v_4$	$v_5$	$v_6$
$v_7$	$v_8$	$v_9$

 $\times$ 

$w_1$	$w_2$
$w_3$	$w_4$

 $=$ 

$G_1$	$G_2$
$G_3$	$G_4$

0  $G_1$


0  $G_2$

0  $G_3$

0  $G_4$

# Pooling/Downsampling Layer

- Commonly combined with convolutional layers
- Applies to each channel separately
- Common implementation for images: *max pooling*,  $2 \times 2$  window, stride of 2 (no parameters to learn)
- No weights to learn in the pooling layer itself
- Shrinks layers leading to fewer weights to learn in subsequent layers
- Subsequent learned features have lower spatial resolution, but higher spatial scope



# Convolutional Neural Networks

- Combination of convolutional, pooling, and fully connected layers
- Backprop extends naturally to CNNs
- Major breakthrough: LeNet5 (Yann LeCun et. al, 1998) for handwritten digit recognition
- CNNs now used in Facebook's face recognition system, self driving cars, and many other object recognition systems

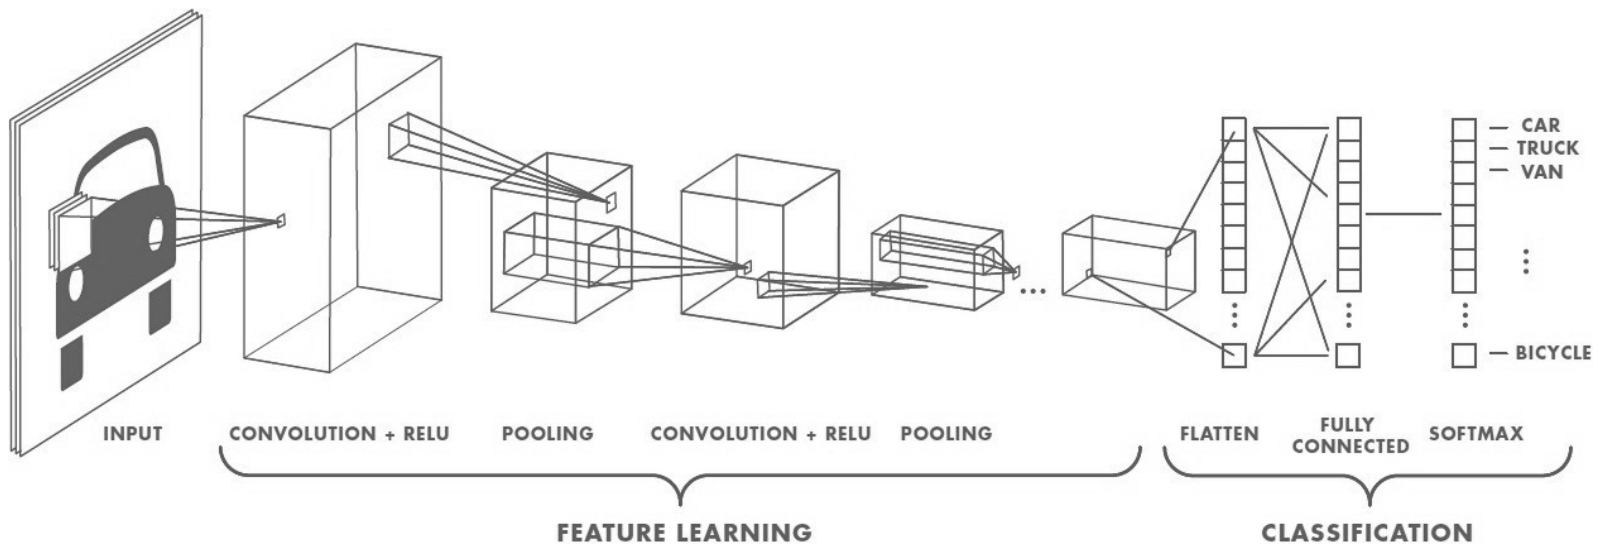


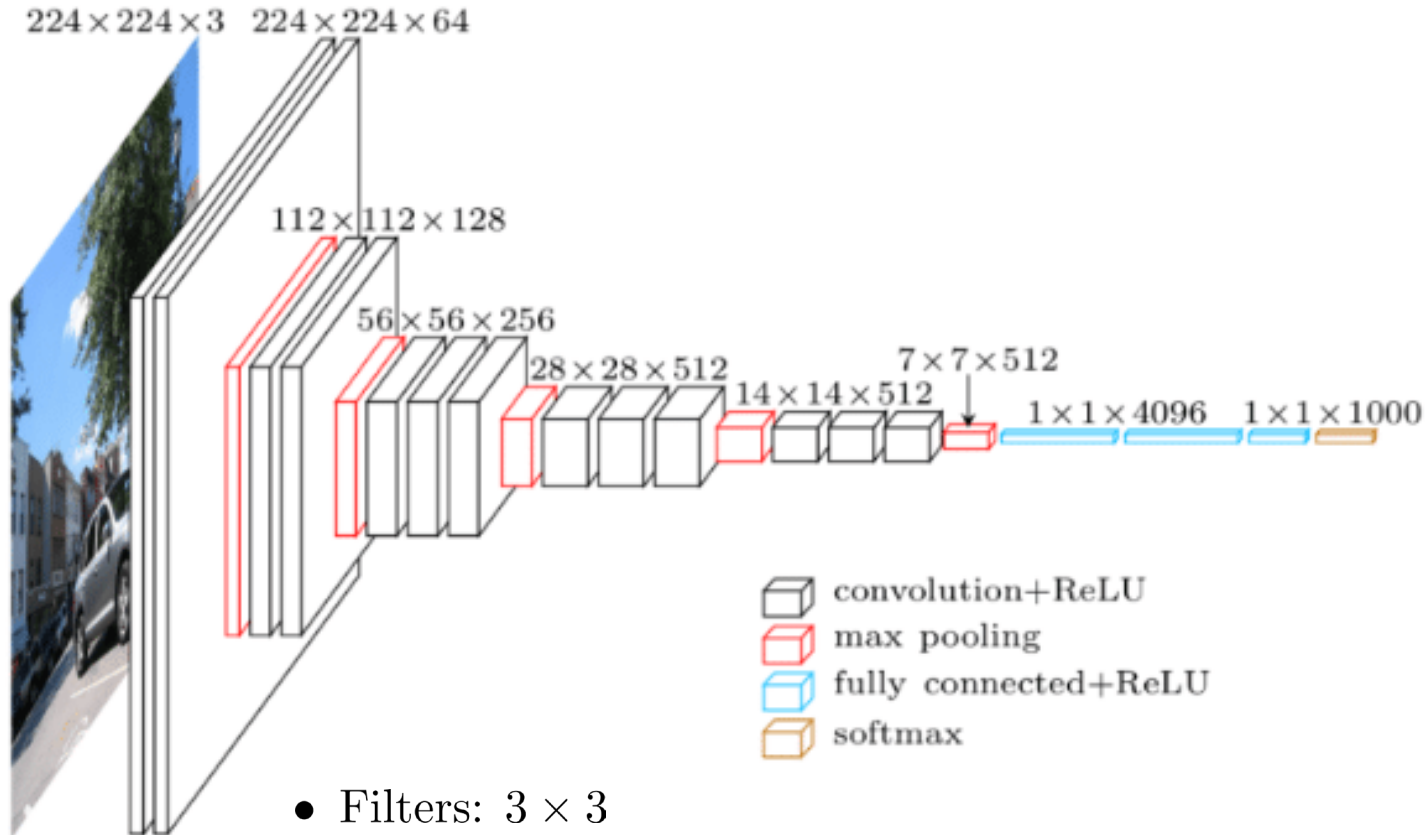
Figure: <https://becominghuman.ai/what-exactly-does-cnn-see-4d436d8e6e52>




# Example: VGG16

VGG16 proposed by K. Simonyan and A. Zisserman, “Very Deep Convolutional Networks for Large-Scale Image Recognition”.

Figure: <https://neurohive.io/en/popular-networks/vgg16/>



# Poll



(A) 0


(B)  $3^2 \times 256$

(C)  $3^2 \times 56^2$

(D)  $3^2 \times 256^2$

(E)  $3^2 \times 56^2 \times 256^2$

# Poll



How many weights connect these two layers?

(A) 0

(B)  $3^2 \times 256$






(C)  $3^2 \times 56^2$

(D)  $3^2 \times 256^2$

(E)  $3^2 \times 56^2 \times 256^2$

# Learned Filters

- Why convolutions (sliding windows)?
  - Fewer weights (as mentioned previously)
  - Salient features are often spatially localized
- Pooling layers lead to “multi-resolution” features
- Below are some filters from VGG16 (trained on a very large image dataset called ImageNet)
- Layers roughly correspond to level of detail.
- Weights are *learned*




# Deep Learning

Enabled by

- Graphics processing units (GPUs): parallel floating-point calculators with 100s-1000s of cores
- Large, public databases such as ImageNet (Fei Fei Li, 2009) which has over 14 million *labeled* examples and 20 thousand classes of objects
- New training strategies
  - Dropout
  - Modifications of SGD (e.g., Adam)
- New architectural elements
  - Residual connections
  - Layer normalization
  - Batch normalization
- Rectified linear units and other activation functions (helps with vanishing gradient problem)


# AlexNet

- The big breakthrough
- Alex Krizhevsky, Ilya Sutskever, and Geoff Hinton (2012)
- Reduced error rate on ImageNet from 26% to 16%
- Used GPUs, dropout, ReLU, which have since become standard



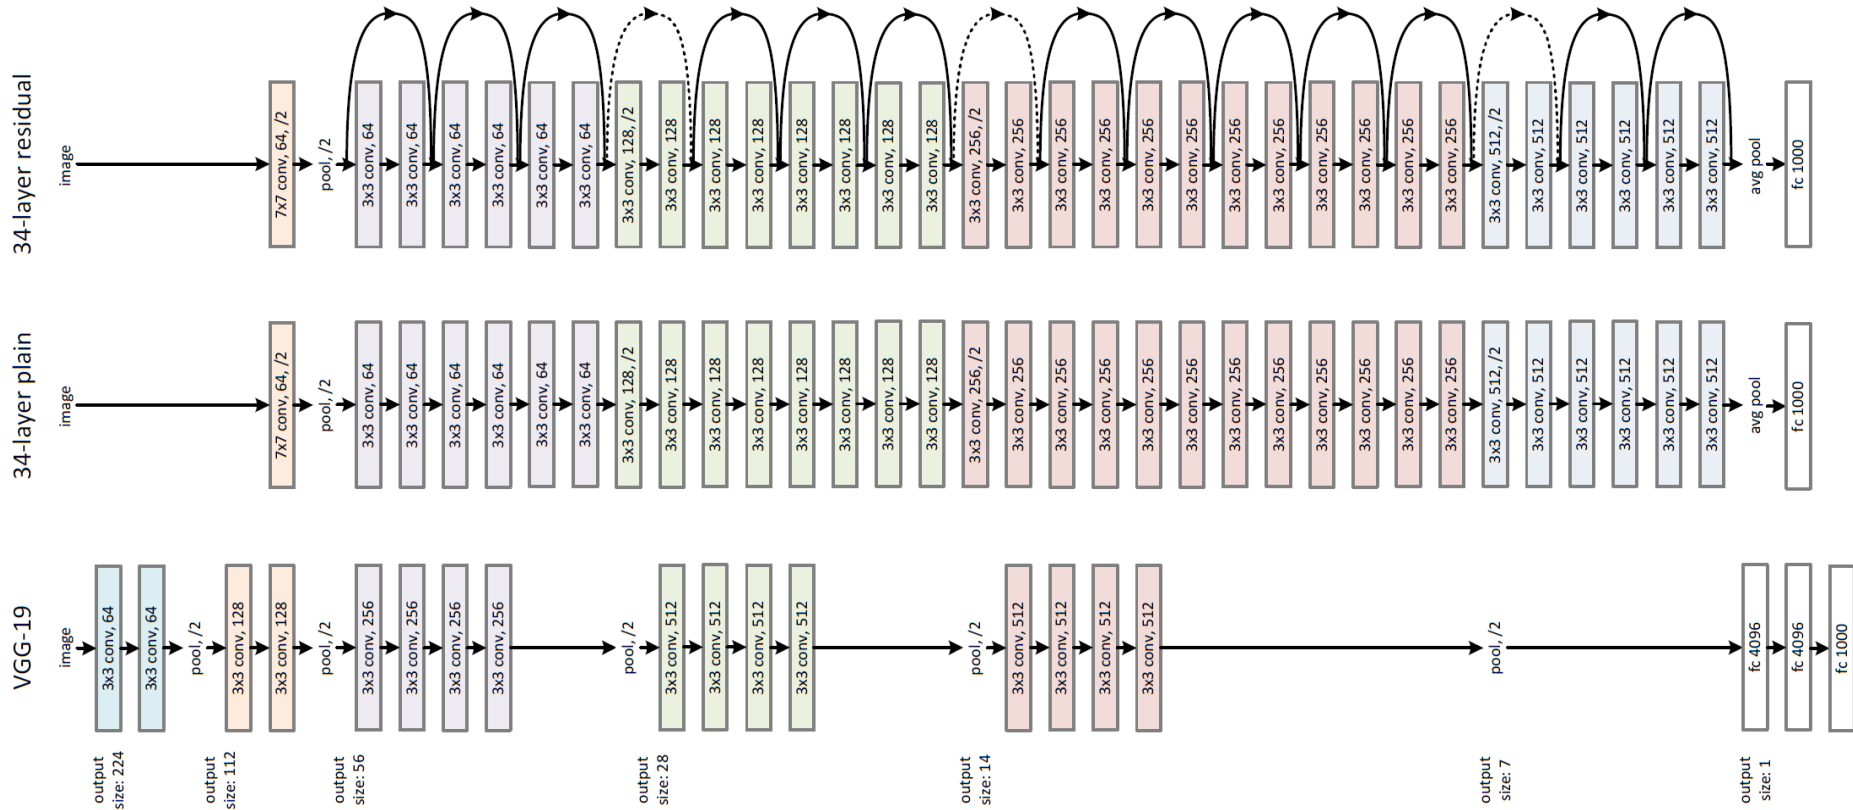
# Residual Networks

- “Residual connections”
- Helps with vanishing gradients; gradient propagates directly back to earlier layers



- Addition is performed before applying activation function

# ResNet




Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun, “Deep Residual Learning for Image Recognition,” CVPR 2015



# Final Thoughts

- CNNs: neural networks for images
- Feedforward but not fully connected; train with backpropagation

# Dilated Convolutions



[https://www.researchgate.net/figure/Dilated-convolution-On-the-left-we-have-the-dilated-convolution-with-dilation-rate-r\\_fig2\\_320195101A](https://www.researchgate.net/figure/Dilated-convolution-On-the-left-we-have-the-dilated-convolution-with-dilation-rate-r_fig2_320195101A)